




1. Bridge 1 and Bridge 3 are in same network. Traffic between the containers connected to bridge 1 and bridge 3 should use L2 Overlay

Here is the topology and its connections:



Here is the list of containers in the hypervisor

```
[sudo] password for ece792:
root@ece792-Standard-PC-i440FX-PIIX-1996:/home/ece792# docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
6f432419fae6	allutil_image	"/bin/bash"	21 hours ago	Up 21 hours		lc2
8015c76fb745	allutil_image	"/bin/bash"	21 hours ago	Up 20 hours		lc1
7119e645b7cb	allutil_image	"/bin/bash"	21 hours ago	Up 21 hours		sc2
1710f5b6ddca	allutil_image	"/bin/bash"	21 hours ago	Up 21 hours		sc1
b80ed10ee269	allutil_image	"/bin/bash"	21 hours ago	Up 21 hours		cs5
2d020b399409	allutil_image	"/bin/bash"	21 hours ago	Up 21 hours		cs4
db273ea1445a	allutil_image	"/bin/bash"	21 hours ago	Up 21 hours		cs3
df6bd55873a3	allutil_image	"/bin/bash"	21 hours ago	Up 21 hours		cs2
c1a6f8c932cd	allutil_image	"/bin/bash"	21 hours ago	Up 21 hours		cs1
14dfdf86ed2d	iserver	"/bin/bash"	23 hours ago	Up 23 hours		CCI
838fa13bcdcf	iserver	"/bin/bash"	23 hours ago	Up 23 hours		CSI
c98146bde2f2	ubuntu	"/bin/bash"	42 hours ago	Up 42 hours		base_image

```
root@ece792-Standard-PC-i440FX-PIIX-1996:/home/ece792#
```

Here is the snapshots of cs1 & cs2 & cs5 containers configs & routes:

CS1-config & Forwarding table :

```
rtt min/avg/max/mdev = 0.256/0.301/0.346/0.045 ms
root@c1a6f8c932cd:/# ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: gre0@NONE: <NOARP> mtu 1476 qdisc noop state DOWN group default qlen 1000
    link/gre 0.0.0.0 brd 0.0.0.0
3: gretap@NONE: <BROADCAST,MULTICAST> mtu 1462 qdisc noop state DOWN group default qlen 1000
    link/ether 00:00:00:00:00:00 brd ff:ff:ff:ff:ff:ff
7: cs1-br1@f8: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
    link/ether ca:89:60:e4:98:80 brd ff:ff:ff:ff:ff:ff link-netnsid 1
    inet 20.0.0.1/24 scope global cs1-br1
        valid_lft forever preferred_lft forever
    inet6 fe80::c889:60ff:fe4:9880/64 scope link
        valid_lft forever preferred_lft forever
458: eth0@f459: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default
    link/ether 02:42:ac:11:00:05 brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet 172.17.0.5/16 scope global eth0
        valid_lft forever preferred_lft forever
    inet6 fe80::42:acff:fe11:5/64 scope link
        valid_lft forever preferred_lft forever
root@c1a6f8c932cd:/# route -n
```

Destination	Gateway	Genmask	Flags	Metric	Ref	Use	Iface
0.0.0.0	172.17.0.1	0.0.0.0	UG	0	0	0	eth0
20.0.0.0	0.0.0.0	255.255.255.0	U	0	0	0	cs1-br1
30.0.0.0	20.0.0.6	255.255.255.0	UG	0	0	0	cs1-br1
172.17.0.0	0.0.0.0	255.255.0.0	U	0	0	0	eth0

```
root@c1a6f8c932cd:/#
```

CS2 config & forwarding table :

```
rtt min/avg/max/mdev = 0.273/0.383/0.493/0.110 ms
root@df6bd55873a3:/# ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: gre0@NONE: <NOARP> mtu 1476 qdisc noop state DOWN group default qlen 1000
    link/gre 0.0.0.0 brd 0.0.0.0
3: gretap@NONE: <BROADCAST,MULTICAST> mtu 1462 qdisc noop state DOWN group default qlen 1000
    link/ether 00:00:00:00:00:00 brd ff:ff:ff:ff:ff:ff
11: cs2-br1@f12: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
    link/ether 66:4a:19:31:4d:16 brd ff:ff:ff:ff:ff:ff link-netnsid 1
    inet 20.0.0.2/24 scope global cs2-br1
        valid_lft forever preferred_lft forever
    inet6 fe80::644a:19ff:fe31:4d16/64 scope link
        valid_lft forever preferred_lft forever
460: eth0@f461: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default
    link/ether 02:42:ac:11:00:06 brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet 172.17.0.6/16 scope global eth0
        valid_lft forever preferred_lft forever
    inet6 fe80::42:acff:fe11:6/64 scope link
        valid_lft forever preferred_lft forever
root@df6bd55873a3:/# route -n
```

Destination	Gateway	Genmask	Flags	Metric	Ref	Use	Iface
0.0.0.0	172.17.0.1	0.0.0.0	UG	0	0	0	eth0
20.0.0.0	0.0.0.0	255.255.255.0	U	0	0	0	cs2-br1
172.17.0.0	0.0.0.0	255.255.0.0	U	0	0	0	eth0

```
root@df6bd55873a3:/#
```

CS3 config & forwarding table :

```
root@db273ea1445a:/# ip route add 20.0.0.0/24 via 30.0.0.3
root@db273ea1445a:/# ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: gre0@NONE: <NOARP> mtu 1476 qdisc noop state DOWN group default qlen 1000
    link/gre 0.0.0.0 brd 0.0.0.0
3: gretap@NONE: <BROADCAST,MULTICAST> mtu 1462 qdisc noop state DOWN group default qlen 1000
    link/ether 00:00:00:00:00:00 brd ff:ff:ff:ff:ff:ff
8: cs3-br@if7: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
    link/ether 52:89:ed:b4:97:30 brd ff:ff:ff:ff:ff:ff link-netnsid 1
    inet 30.0.0.1/24 scope global cs3-br
        valid_lft forever preferred_lft forever
    inet6 fe80::5009:edff:feb4:9730/64 scope link
        valid_lft forever preferred_lft forever
462: eth0@if463: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default
    link/ether 02:42:ac:11:00:07 brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet 172.17.0.7/16 scope global eth0
        valid_lft forever preferred_lft forever
    inet6 fe80::42:acff:fe11:7/64 scope link
        valid_lft forever preferred_lft forever
root@db273ea1445a:/# route -n
Kernel IP routing table
Destination Gateway Genmask Flags Metric Ref Use Iface
0.0.0.0 172.17.0.1 0.0.0.0 UG 0 0 0 eth0
20.0.0.0 30.0.0.3 255.255.255.0 UG 0 0 0 cs3-br
30.0.0.0 0.0.0.0 255.255.255.0 U 0 0 0 cs3-br
172.17.0.0 0.0.0.0 255.255.0.0 U 0 0 0 eth0
root@db273ea1445a:/#
```

CS4 config & forwarding table:

```
rtt min/avg/max/mdev = 0.086/0.086/0.087/0.009 ms
root@d22bb398409:/# ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: gre0@NONE: <NOARP> mtu 1476 qdisc noop state DOWN group default qlen 1000
    link/gre 0.0.0.0 brd 0.0.0.0
3: gretap@NONE: <BROADCAST,MULTICAST> mtu 1462 qdisc noop state DOWN group default qlen 1000
    link/ether 00:00:00:00:00:00 brd ff:ff:ff:ff:ff:ff
9: cs4-br@if10: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
    link/ether aa:13:f7:9e:bdc8 brd ff:ff:ff:ff:ff:ff link-netnsid 1
    inet 30.0.0.2/24 scope global cs4-br
        valid_lft forever preferred_lft forever
    inet6 fe80::ac13:f7ff:fe0e:bdc8/64 scope link
        valid_lft forever preferred_lft forever
464: eth0@if465: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default
    link/ether 02:42:ac:11:00:08 brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet 172.17.0.8/16 scope global eth0
        valid_lft forever preferred_lft forever
    inet6 fe80::42:acff:fe11:8/64 scope link
        valid_lft forever preferred_lft forever
root@d22bb398409:/# route -n
Kernel IP routing table
Destination Gateway Genmask Flags Metric Ref Use Iface
0.0.0.0 172.17.0.1 0.0.0.0 UG 0 0 0 eth0
20.0.0.0 30.0.0.3 255.255.255.0 UG 0 0 0 cs4-br
30.0.0.0 0.0.0.0 255.255.255.0 U 0 0 0 cs4-br
172.17.0.0 0.0.0.0 255.255.0.0 U 0 0 0 eth0
root@d22bb398409:/#
```

CS5 config & forwarding table :

```
Last login: Thu Nov 29 18:33:05 2018 from 192.168.122.1
ece792@ece792-Standard-PC-1440FX-PIIX-1996:~$ sudo su
[sudo] password for ece792:
root@ece792-Standard-PC-1440FX-PIIX-1996:/home/ece792# docker exec -it --privileged cs5 bash
root@b80ed10ee269:/# ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: gre0@NONE: <NOARP> mtu 1476 qdisc noop state DOWN group default qlen 1000
    link/gre 0.0.0.0 brd 0.0.0.0
3: gretap@NONE: <BROADCAST,MULTICAST> mtu 1462 qdisc noop state DOWN group default qlen 1000
    link/ether 00:00:00:00:00:00 brd ff:ff:ff:ff:ff:ff
5: cs5-br3@if6: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
    link/ether ea:36:90:91:bd:87 brd ff:ff:ff:ff:ff:ff link-netnsid 1
    inet 20.0.0.3/24 scope global cs5-br3
        valid_lft forever preferred_lft forever
    inet6 fe80::e836:90ff:fe91:bd87/64 scope link
        valid_lft forever preferred_lft forever
466: eth0@if467: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default
    link/ether 02:42:ac:11:00:09 brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet 172.17.0.9/16 scope global eth0
        valid_lft forever preferred_lft forever
    inet6 fe80::42:acff:fe11:9/64 scope link
        valid_lft forever preferred_lft forever
root@b80ed10ee269:/# route -n
Kernel IP routing table
Destination Gateway Genmask Flags Metric Ref Use Iface
0.0.0.0 172.17.0.1 0.0.0.0 UG 0 0 0 eth0
20.0.0.0 0.0.0.0 255.255.255.0 U 0 0 0 cs5-br3
172.17.0.0 0.0.0.0 255.255.0.0 U 0 0 0 eth0
root@b80ed10ee269:/#
```

Here is the config of con_ns1 namespace which has the bridge (bridge1) connected to cs1 and cs2

```
1: lo: <LOOPBACK> mtu 65536 qdisc noop state DOWN group default qlen 1000
   link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: gre0@NONE: <NOARP> mtu 1476 qdisc noop state DOWN group default qlen 1000
   link/gre 0.0.0.0 brd 0.0.0.0
3: gretap0@NONE: <BROADCAST,MULTICAST> mtu 1462 qdisc noop state DOWN group default qlen 1000
   link/ether 00:00:00:00:00:00 brd ff:ff:ff:ff:ff:ff
5: bridge1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1450 qdisc noqueue state UP group default qlen 1000
   link/ether 3e:aa:9f:ea:ef:80 brd ff:ff:ff:ff:ff:ff
   inet6 fe80::7813:aff:fe4:f0f8/64 scope link
       valid_lft forever preferred_lft forever
8: br1_cs1@if7: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue master bridge1 state UP group default qlen 1000
   link/ether 7a:13:0a:f4:f0:f8 brd ff:ff:ff:ff:ff:ff link-netnsid 0
   inet6 fe80::7013:aff:fe4:f0f8/64 scope link
       valid_lft forever preferred_lft forever
12: br1_cs2@if11: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue master bridge1 state UP group default qlen 1000
   link/ether b6:ea:1f:d5:43:1a brd ff:ff:ff:ff:ff:ff link-netnsid 1
   inet6 fe80::b4ea:1fff:fed5:431a/64 scope link
       valid_lft forever preferred_lft forever
14: br1_lc1@if13: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
   link/ether be:c8:ee:d3:1f brd ff:ff:ff:ff:ff:ff link-netnsid 2
   inet 21.0.0.1/24 scope global br1_lc1
       valid_lft forever preferred_lft forever
   inet6 fe80::bcc8:eeff:feda:d31f/64 scope link
       valid_lft forever preferred_lft forever
15: vxlan-cont: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1450 qdisc noqueue master bridge1 state UNKNOWN group default qlen 1000
   link/ether 3e:aa:9f:ea:ef:80 brd ff:ff:ff:ff:ff:ff
   inet6 fe80::3caa:9fff:feea:ef80/64 scope link
       valid_lft forever preferred_lft forever
root@ece792-Standard-PC-i440FX-PIIX-1996:/home/ece792# route -n
Kernel IP routing table
Destination    Gateway         Genmask         Flags Metric Ref    Use Iface
0.0.0.0        21.0.0.2       0.0.0.0         UG    0     0      0 br1_lc1
21.0.0.0       0.0.0.0        255.255.255.0   U     0     0      0 br1_lc1
root@ece792-Standard-PC-i440FX-PIIX-1996:/home/ece792#
```

Here is the config of con_ns3 which has the bridge3(br3) connected to cs5 container:

```
root@ece792-Standard-PC-i440FX-PIIX-1996:/home/ece792# ip netns exec con_ns3 bash
root@ece792-Standard-PC-i440FX-PIIX-1996:/home/ece792# ip a
1: lo: <LOOPBACK> mtu 65536 qdisc noop state DOWN group default qlen 1000
   link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: gre0@NONE: <NOARP> mtu 1476 qdisc noop state DOWN group default qlen 1000
   link/gre 0.0.0.0 brd 0.0.0.0
3: gretap0@NONE: <BROADCAST,MULTICAST> mtu 1462 qdisc noop state DOWN group default qlen 1000
   link/ether 00:00:00:00:00:00 brd ff:ff:ff:ff:ff:ff
4: bridge3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1450 qdisc noqueue state UP group default qlen 1000
   link/ether 32:c6:8c:56:1d:70 brd ff:ff:ff:ff:ff:ff
   inet6 fe80::c81c:18ff:fe72:d01/64 scope link
       valid_lft forever preferred_lft forever
6: br3_cs5@if5: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue master bridge3 state UP group default qlen 1000
   link/ether 32:c6:8c:56:1d:70 brd ff:ff:ff:ff:ff:ff link-netnsid 0
   inet6 fe80::30c6:8cff:fe56:1d70/64 scope link
       valid_lft forever preferred_lft forever
9: br3_lc2@if10: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
   link/ether 82:90:ac:22:ab:05 brd ff:ff:ff:ff:ff:ff link-netnsid 1
   inet 23.0.0.1/24 scope global br3_lc2
       valid_lft forever preferred_lft forever
   inet6 fe80::8090:acff:fe22:ab05/64 scope link
       valid_lft forever preferred_lft forever
13: vxlan-cont: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1450 qdisc noqueue master bridge3 state UNKNOWN group default qlen 1000
   link/ether 92:94:d0:1f:7e:44 brd ff:ff:ff:ff:ff:ff
   inet6 fe80::9094:d0ff:fef1:7e44/64 scope link
       valid_lft forever preferred_lft forever
root@ece792-Standard-PC-i440FX-PIIX-1996:/home/ece792# route -n
Kernel IP routing table
Destination    Gateway         Genmask         Flags Metric Ref    Use Iface
0.0.0.0        23.0.0.2       0.0.0.0         UG    0     0      0 br3_lc2
23.0.0.0       0.0.0.0        255.255.255.0   U     0     0      0 br3_lc2
root@ece792-Standard-PC-i440FX-PIIX-1996:/home/ece792#
```

Here are the configs of lc1 & lc2 containers :

Lc1 config (leaf container1)


```

link/ether 00:00:00:00:00:00 brd ff:ff:ff:ff:ff:ff
4: gretun-cont@NONE: <POINTOPOINT,NOARP,UP,LOWER_UP> mtu 1476 qdisc noqueue state UNKNOWN group default qlen 1000
link/gre 26.0.0.2 peer 28.0.0.2
inet6 fe80::200:5efe:1a00:2/64 scope link
    valid_lft forever preferred_lft forever
12: lc1_lc2@if11: <BROADCAST,MULTICAST> mtu 1500 qdisc noqueue state DOWN group default qlen 1000
link/ether 8a:97:01:83:e1:33 brd ff:ff:ff:ff:ff:ff link-netnsid 5
inet 22.0.0.2/24 scope global lc1_lc2
    valid_lft forever preferred_lft forever
13: lc1_br1@if14: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
link/ether 32:fb:c6:aa:73:50 brd ff:ff:ff:ff:ff:ff link-netnsid 3
inet 21.0.0.2/24 scope global lc1_br1
    valid_lft forever preferred_lft forever
inet6 fe80::30fb:c6ff:feaa:7350/64 scope link
    valid_lft forever preferred_lft forever
484: eth0@if485: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default
link/ether 02:42:ac:11:00:0c brd ff:ff:ff:ff:ff:ff link-netnsid 0
inet 172.17.0.12/16 scope global eth0
    valid_lft forever preferred_lft forever
inet6 fe80::42:acff:fe11:c/64 scope link
    valid_lft forever preferred_lft forever
487: lc1_sc1@if486: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
link/ether 02:f8:d0:dc:03:b4 brd ff:ff:ff:ff:ff:ff link-netnsid 1
inet 25.0.0.2/24 scope global lc1_sc1
    valid_lft forever preferred_lft forever
inet6 fe80::f8:d0ff:fedc:3b4/64 scope link
    valid_lft forever preferred_lft forever
489: lc1_sc2@if488: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
link/ether 72:de:43:b4:d1:64 brd ff:ff:ff:ff:ff:ff link-netnsid 2
inet 26.0.0.2/24 scope global lc1_sc2
    valid_lft forever preferred_lft forever
inet6 fe80::70de:43ff:feb4:d164/64 scope link
    valid_lft forever preferred_lft forever
root@8015c76fb745:/#

```

Forwarding table of lc1:

```

Kernel IP routing table

```

Destination	Gateway	Genmask	Flags	Metric	Ref	Use	Iface
0.0.0.0	172.17.0.1	0.0.0.0	UG	0	0	0	eth0
20.0.0.0	0.0.0.0	255.255.255.0	U	0	0	0	lc1-bridge1
21.0.0.0	0.0.0.0	255.255.255.0	U	0	0	0	lc1_br1
23.0.0.0	0.0.0.0	255.255.255.0	U	0	0	0	gretun-cont
25.0.0.0	0.0.0.0	255.255.255.0	U	0	0	0	lc1_sc1
26.0.0.0	0.0.0.0	255.255.255.0	U	0	0	0	lc1_sc2
27.0.0.0	25.0.0.1	255.255.255.0	UG	0	0	0	lc1_sc1
28.0.0.0	26.0.0.1	255.255.255.0	UG	0	0	0	lc1_sc2
30.0.0.0	0.0.0.0	255.255.255.0	U	0	0	0	gretun-cont
41.0.0.0	0.0.0.0	255.255.255.0	U	0	0	0	lc1-csX-ns
43.0.0.0	26.0.0.1	255.255.255.0	UG	0	0	0	lc1_sc2
100.0.0.0	0.0.0.0	255.255.255.0	U	0	0	0	lc1-csX1-br
101.0.0.0	0.0.0.0	255.255.255.0	U	0	0	0	gretun-csX1
110.0.0.0	0.0.0.0	255.255.255.0	U	0	0	0	lc1-csA-br
111.0.0.0	26.0.0.1	255.255.255.0	UG	0	0	0	lc1_sc2
172.17.0.0	0.0.0.0	255.255.0.0	U	0	0	0	eth0

```

root@8015c76fb745:/#

```

LC2 config (lc2) config :

```

6: gretun-cont@NONE: <POINTOPOINT,NOARP,UP,LOWER_UP> mtu 1476 qdisc noqueue state UNKNOWN group default qlen 1000
link/gre 28.0.0.2 peer 26.0.0.2
inet6 fe80::200:5efe:1c00:2/64 scope link
    valid_lft forever preferred_lft forever
10: lc2_br3@if9: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
link/ether ea:31:a6:3c:b0:32 brd ff:ff:ff:ff:ff:ff link-netnsid 3
inet 23.0.0.2/24 scope global lc2_br3
    valid_lft forever preferred_lft forever
inet6 fe80::e831:a6ff:fe3c:b032/64 scope link
    valid_lft forever preferred_lft forever
11: lc2_lc1@if12: <BROADCAST,MULTICAST> mtu 1500 qdisc noqueue state DOWN group default qlen 1000
link/ether 5e:cd:16:70:e7:08 brd ff:ff:ff:ff:ff:ff link-netnsid 4
inet 22.0.0.3/24 scope global lc2_lc1
    valid_lft forever preferred_lft forever
474: eth0@if475: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default
link/ether 02:42:ac:11:00:0d brd ff:ff:ff:ff:ff:ff link-netnsid 0
inet 172.17.0.13/16 scope global eth0
    valid_lft forever preferred_lft forever
inet6 fe80::42:acff:fe11:d/64 scope link
    valid_lft forever preferred_lft forever
481: lc2_sc1@if480: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
link/ether 52:b0:b1:40:4a:a0 brd ff:ff:ff:ff:ff:ff link-netnsid 1
inet 27.0.0.2/24 scope global lc2_sc1
    valid_lft forever preferred_lft forever
inet6 fe80::50b0:b1ff:fe40:4aa0/64 scope link
    valid_lft forever preferred_lft forever
483: lc2_sc2@if482: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
link/ether 2a:74:8f:95:32:78 brd ff:ff:ff:ff:ff:ff link-netnsid 2
inet 28.0.0.2/24 scope global lc2_sc2
    valid_lft forever preferred_lft forever
inet6 fe80::2874:8fff:fe95:3278/64 scope link
    valid_lft forever preferred_lft forever
root@6f432419fae6:/#

```

Forwarding table of lc2:

```
root@6f432419fae6:/# route -n
Kernel IP routing table
Destination      Gateway         Genmask         Flags Metric Ref    Use Iface
0.0.0.0          172.17.0.1     0.0.0.0         UG    0      0          0 eth0
20.0.0.0         0.0.0.0        255.255.255.0   U      0      0          0 gretun-cont
21.0.0.0         0.0.0.0        255.255.255.0   U      0      0          0 gretun-cont
23.0.0.0         0.0.0.0        255.255.255.0   U      0      0          0 lc2_br3
25.0.0.0         27.0.0.1      255.255.255.0   UG    0      0          0 lc2_sc1
26.0.0.0         28.0.0.1      255.255.255.0   UG    0      0          0 lc2_sc2
27.0.0.0         0.0.0.0        255.255.255.0   U      0      0          0 lc2_sc1
28.0.0.0         0.0.0.0        255.255.255.0   U      0      0          0 lc2_sc2
30.0.0.0         0.0.0.0        255.255.255.0   U      0      0          0 lc2-br2
41.0.0.0         28.0.0.1      255.255.255.0   UG    0      0          0 lc2_sc2
43.0.0.0         0.0.0.0        255.255.255.0   U      0      0          0 lc2-csY-ns
100.0.0.0        0.0.0.0        255.255.255.0   U      0      0          0 gretun-csY1
101.0.0.0        0.0.0.0        255.255.255.0   U      0      0          0 lc2-csY1-br
110.0.0.0        28.0.0.1      255.255.255.0   UG    0      0          0 lc2_sc2
111.0.0.0        0.0.0.0        255.255.255.0   U      0      0          0 lc2-csB-br
172.17.0.0       0.0.0.0        255.255.0.0     U      0      0          0 eth0
root@6f432419fae6:/#
```

sc1 config & forward table:

```
inet 127.0.0.1/8 scope host lo
    valid_lft forever preferred_lft forever
inet6 ::1/128 scope host
    valid_lft forever preferred_lft forever
2: gre0@NONE: <NOARP> mtu 1476 qdisc noop state DOWN group default qlen 1000
    link/gre 0.0.0.0 brd 0.0.0.0
3: gretap@NONE: <BROADCAST,MULTICAST> mtu 1462 qdisc noop state DOWN group default qlen 1000
    link/ether 00:00:00:00:00:00 brd ff:ff:ff:ff:ff:ff
468: eth0@if469: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default
    link/ether 02:42:ac:11:00:0a brd ff:ff:ff:ff:ff:ff link-netnsid 0
    inet 172.17.0.10/16 scope global eth0
        valid_lft forever preferred_lft forever
    inet6 fe80::42:acff:fe11:a/64 scope link
        valid_lft forever preferred_lft forever
480: sc1_lc2@if481: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
    link/ether 6e:6f:74:a1:53:cc brd ff:ff:ff:ff:ff:ff link-netnsid 2
    inet 27.0.0.1/24 scope global sc1_lc2
        valid_lft forever preferred_lft forever
    inet6 fe80::6c6f:74ff:feal:53cc/64 scope link
        valid_lft forever preferred_lft forever
486: sc1_lc1@if487: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
    link/ether 72:45:bc:3e:1e:a1 brd ff:ff:ff:ff:ff:ff link-netnsid 1
    inet 25.0.0.1/24 scope global sc1_lc1
        valid_lft forever preferred_lft forever
    inet6 fe80::7045:bcff:fe3e:1eal/64 scope link
        valid_lft forever preferred_lft forever
root@1710f5b6ddca:/# route -n
Kernel IP routing table
Destination      Gateway         Genmask         Flags Metric Ref    Use Iface
0.0.0.0          172.17.0.1     0.0.0.0         UG    0      0          0 eth0
25.0.0.0         0.0.0.0        255.255.255.0   U      0      0          0 sc1_lc1
27.0.0.0         0.0.0.0        255.255.255.0   U      0      0          0 sc1_lc2
172.17.0.0       0.0.0.0        255.255.0.0     U      0      0          0 eth0
root@1710f5b6ddca:/#
```

Sc2 config & forwarding table:

```

inet 127.0.0.1/8 scope host lo
    valid_lft forever preferred_lft forever
inet6 ::1/128 scope host
    valid_lft forever preferred_lft forever
2: gre0@NONE: <NOARP> mtu 1476 qdisc noop state DOWN group default qlen 1000
    link/gre 0.0.0.0 brd 0.0.0.0
3: gretap@NONE: <BROADCAST,MULTICAST> mtu 1462 qdisc noop state DOWN group default qlen 1000
    link/ether 00:00:00:00:00:00 brd ff:ff:ff:ff:ff:ff
470: eth0@if471: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default
    link/ether 02:42:ac:11:00:0b brd ff:ff:ff:ff:ff:ff link-netnsid 0
inet 172.17.0.11/16 scope global eth0
    valid_lft forever preferred_lft forever
inet6 fe80::42:acff:fe11:b/64 scope link
    valid_lft forever preferred_lft forever
482: sc2_lc2@if483: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
    link/ether ca:81:a7:25:97:67 brd ff:ff:ff:ff:ff:ff link-netnsid 2
inet 28.0.0.1/24 scope global sc2_lc2
    valid_lft forever preferred_lft forever
inet6 fe80::c881:a7ff:fe25:9767/64 scope link
    valid_lft forever preferred_lft forever
488: sc2_lc1@if489: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
    link/ether 06:15:d5:2e:26:61 brd ff:ff:ff:ff:ff:ff link-netnsid 1
inet 26.0.0.1/24 scope global sc2_lc1
    valid_lft forever preferred_lft forever
inet6 fe80::415:d5ff:fe2e:2661/64 scope link
    valid_lft forever preferred_lft forever
root@7119e645b7cb:/# route -n
Kernel IP routing table

```

Destination	Gateway	Genmask	Flags	Metric	Ref	Use	Iface
0.0.0.0	172.17.0.1	0.0.0.0	UG	0	0	0	eth0
26.0.0.0	0.0.0.0	255.255.255.0	U	0	0	0	sc2_lc1
28.0.0.0	0.0.0.0	255.255.255.0	U	0	0	0	sc2_lc2
172.17.0.0	0.0.0.0	255.255.0.0	U	0	0	0	eth0

```

root@7119e645b7cb:/#

```

Vxlan configuration in con_ns1 namespace:

```

root@ece792-Standard-PC-i440FX-PIIX-1996:/home/ece792# ip -d link vxlan-cont
Command "vxlan-cont" is unknown, try "ip link help".
root@ece792-Standard-PC-i440FX-PIIX-1996:/home/ece792# ip -d link show vxlan-cont
15: vxlan-cont: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1450 qdisc noqueue master bridge1 state UNKNOWN mode DEFAULT group default qlen 1000
    link/ether 3e:aa:9f:ea:ef:80 brd ff:ff:ff:ff:ff:ff promiscuity 1
    vxlan id 30 remote 23.0.0.1 dev br1_lc1 srcport 0 0 dstport 4900 ageing 300
    bridge_slave state forwarding priority 32 cost 100 hairpin off guard off root_block off fastleave off learning on flood on addrngmode eui64
root@ece792-Standard-PC-i440FX-PIIX-1996:/home/ece792#

```

Vxlan configuration in con_ns3 namespace :

```

root@ece792-Standard-PC-i440FX-PIIX-1996:/home/ece792# ip -d link show vxlan-cont
13: vxlan-cont: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1450 qdisc noqueue master bridge3 state UNKNOWN mode DEFAULT group default qlen 1000
    link/ether 92:94:d0:1f:7e:44 brd ff:ff:ff:ff:ff:ff promiscuity 1
    vxlan id 30 remote 21.0.0.1 dev br3_lc2 srcport 0 0 dstport 4900 ageing 300
    bridge_slave state forwarding priority 32 cost 100 hairpin off guard off root_block off fastleave off learning on flood on addrngmode eui64
root@ece792-Standard-PC-i440FX-PIIX-1996:/home/ece792#

```

Gre tunnel configuration in lc1:

```

root@8015c76fb745:/# ip tunnel show
gretun-cont: gre/ip remote 28.0.0.2 local 26.0.0.2 ttl inherit
gre0: gre/ip remote any local any ttl inherit nopmtudisc
root@8015c76fb745:/#

```

Gre tunnel configuration in lc2:

```

root@6f432419fae6:/# ip tunnel show
gretun-cont: gre/ip remote 26.0.0.2 local 28.0.0.2 ttl inherit
gre0: gre/ip remote any local any ttl inherit nopmtudisc
root@6f432419fae6:/#

```

Packet capture at cs1 (20.0.0.1) connected to lc1 via bridge1 to cs5 (20.0.0.3) which is connected to bridge3 to lc2:

Packet capture at cs1:

```
root@cla6f8c932cd:/# ping 20.0.0.3
PING 20.0.0.3 (20.0.0.3) 56(84) bytes of data.
64 bytes from 20.0.0.3: icmp_seq=1 ttl=64 time=0.271 ms
64 bytes from 20.0.0.3: icmp_seq=2 ttl=64 time=0.275 ms
64 bytes from 20.0.0.3: icmp_seq=3 ttl=64 time=0.242 ms
64 bytes from 20.0.0.3: icmp_seq=4 ttl=64 time=0.280 ms
64 bytes from 20.0.0.3: icmp_seq=5 ttl=64 time=0.224 ms
64 bytes from 20.0.0.3: icmp_seq=6 ttl=64 time=0.264 ms
64 bytes from 20.0.0.3: icmp_seq=7 ttl=64 time=0.195 ms
64 bytes from 20.0.0.3: icmp_seq=8 ttl=64 time=0.192 ms
64 bytes from 20.0.0.3: icmp_seq=9 ttl=64 time=0.199 ms
64 bytes from 20.0.0.3: icmp_seq=10 ttl=64 time=0.269 ms
64 bytes from 20.0.0.3: icmp_seq=11 ttl=64 time=0.194 ms
64 bytes from 20.0.0.3: icmp_seq=12 ttl=64 time=0.189 ms
^C
--- 20.0.0.3 ping statistics ---
12 packets transmitted, 12 received, 0% packet loss, time 11249ms
rtt min/avg/max/mdev = 0.189/0.232/0.280/0.041 ms
root@cla6f8c932cd:/# ping 20.0.0.2
PING 20.0.0.2 (20.0.0.2) 56(84) bytes of data.
64 bytes from 20.0.0.2: icmp_seq=1 ttl=64 time=0.172 ms
64 bytes from 20.0.0.2: icmp_seq=2 ttl=64 time=0.084 ms
64 bytes from 20.0.0.2: icmp_seq=3 ttl=64 time=0.084 ms
64 bytes from 20.0.0.2: icmp_seq=4 ttl=64 time=0.090 ms
64 bytes from 20.0.0.2: icmp_seq=5 ttl=64 time=0.094 ms
64 bytes from 20.0.0.2: icmp_seq=6 ttl=64 time=0.108 ms
64 bytes from 20.0.0.2: icmp_seq=7 ttl=64 time=0.108 ms
64 bytes from 20.0.0.2: icmp_seq=8 ttl=64 time=0.087 ms
64 bytes from 20.0.0.2: icmp_seq=9 ttl=64 time=0.083 ms
^C
--- 20.0.0.2 ping statistics ---
9 packets transmitted, 9 received, 0% packet loss, time 8198ms
rtt min/avg/max/mdev = 0.083/0.101/0.172/0.027 ms
root@cla6f8c932cd:/#
```

Packet capture at **lc1-br1**(after vxlan encapsulation at lc1 entry interface):

VXLAN tunnel encapsulation at con_ns1 namespace:

local: 21.0.0.1

remote :23.0.0.1

```
tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
listening on lc1_br1, link-type EN10MB (Ethernet), capture size 262144 bytes
23:29:04.283293 IP 21.0.0.1.40809 > 23.0.0.1.4900: UDP, length 106
23:29:04.283438 IP 23.0.0.1.40809 > 21.0.0.1.4900: UDP, length 106
23:29:05.313148 IP 21.0.0.1.40809 > 23.0.0.1.4900: UDP, length 106
23:29:05.313335 IP 23.0.0.1.40809 > 21.0.0.1.4900: UDP, length 106
23:29:06.337191 IP 21.0.0.1.40809 > 23.0.0.1.4900: UDP, length 106
23:29:06.340126 IP 23.0.0.1.40809 > 21.0.0.1.4900: UDP, length 106
23:29:07.338395 IP 21.0.0.1.40809 > 23.0.0.1.4900: UDP, length 106
23:29:07.338592 IP 23.0.0.1.40809 > 21.0.0.1.4900: UDP, length 106
23:29:08.353429 IP 21.0.0.1.40809 > 23.0.0.1.4900: UDP, length 106
23:29:08.353566 IP 23.0.0.1.40809 > 21.0.0.1.4900: UDP, length 106
23:29:09.345001 ARP, Request who-has 21.0.0.1 tell 21.0.0.2, length 28
23:29:09.345071 ARP, Request who-has 21.0.0.2 tell 21.0.0.1, length 28
23:29:09.345259 ARP, Reply 21.0.0.2 is-at 32:fb:c6:aa:73:50, length 28
23:29:09.345166 ARP, Reply 21.0.0.1 is-at be:c8:ee:da:d3:1f, length 28
23:29:09.345306 IP 21.0.0.1.35691 > 23.0.0.1.4900: UDP, length 50
23:29:09.345415 IP 23.0.0.1.35691 > 21.0.0.1.4900: UDP, length 50
23:29:09.345477 IP 21.0.0.1.35691 > 23.0.0.1.4900: UDP, length 50
23:29:09.377116 IP 21.0.0.1.40809 > 23.0.0.1.4900: UDP, length 106
23:29:09.377238 IP 23.0.0.1.40809 > 21.0.0.1.4900: UDP, length 106
23:29:10.401201 IP 21.0.0.1.40809 > 23.0.0.1.4900: UDP, length 106
23:29:10.401417 IP 23.0.0.1.40809 > 21.0.0.1.4900: UDP, length 106
23:29:11.425124 IP 21.0.0.1.40809 > 23.0.0.1.4900: UDP, length 106
23:29:11.425262 IP 23.0.0.1.40809 > 21.0.0.1.4900: UDP, length 106
23:29:12.449126 IP 21.0.0.1.40809 > 23.0.0.1.4900: UDP, length 106
23:29:12.449266 IP 23.0.0.1.40809 > 21.0.0.1.4900: UDP, length 106
^C
25 packets captured
26 packets received by filter
1 packet dropped by kernel
```



Packet capture at sc2_lc1 interface (at sc2 ----lc1) link after the GRE encapsulation at GRE tunnel in lc1

GRE tunnel at lc1:

Local ip :26.0.0.2

remote ip: 28.0.0.2

```
root@7119e645b/cb:/# tcpdump -i sc2_lc1 -nn
tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
listening on sc2_lc1, link-type EN10MB (Ethernet), capture size 262144 bytes
23:18:47.077109 IP 26.0.0.2 > 28.0.0.2: GREv0, length 138: IP 21.0.0.1.40809 > 23.0.0.1.4900: UDP, length 106
23:18:47.077202 IP 28.0.0.2 > 26.0.0.2: GREv0, length 138: IP 23.0.0.1.40809 > 21.0.0.1.4900: UDP, length 106
23:18:48.097184 IP 26.0.0.2 > 28.0.0.2: GREv0, length 138: IP 21.0.0.1.40809 > 23.0.0.1.4900: UDP, length 106
23:18:48.097325 IP 28.0.0.2 > 26.0.0.2: GREv0, length 138: IP 23.0.0.1.40809 > 21.0.0.1.4900: UDP, length 106
23:18:49.121149 IP 26.0.0.2 > 28.0.0.2: GREv0, length 138: IP 21.0.0.1.40809 > 23.0.0.1.4900: UDP, length 106
23:18:49.121413 IP 28.0.0.2 > 26.0.0.2: GREv0, length 138: IP 23.0.0.1.40809 > 21.0.0.1.4900: UDP, length 106
23:18:50.145131 IP 26.0.0.2 > 28.0.0.2: GREv0, length 138: IP 21.0.0.1.40809 > 23.0.0.1.4900: UDP, length 106
23:18:50.145287 IP 28.0.0.2 > 26.0.0.2: GREv0, length 138: IP 23.0.0.1.40809 > 21.0.0.1.4900: UDP, length 106
23:18:51.169152 IP 26.0.0.2 > 28.0.0.2: GREv0, length 138: IP 21.0.0.1.40809 > 23.0.0.1.4900: UDP, length 106
23:18:51.169301 IP 28.0.0.2 > 26.0.0.2: GREv0, length 138: IP 23.0.0.1.40809 > 21.0.0.1.4900: UDP, length 106
23:18:52.193132 IP 26.0.0.2 > 28.0.0.2: GREv0, length 138: IP 21.0.0.1.40809 > 23.0.0.1.4900: UDP, length 106
23:18:52.193310 IP 28.0.0.2 > 26.0.0.2: GREv0, length 138: IP 23.0.0.1.40809 > 21.0.0.1.4900: UDP, length 106
23:18:53.217179 IP 26.0.0.2 > 28.0.0.2: GREv0, length 138: IP 21.0.0.1.40809 > 23.0.0.1.4900: UDP, length 106
23:18:53.217421 IP 28.0.0.2 > 26.0.0.2: GREv0, length 138: IP 23.0.0.1.40809 > 21.0.0.1.4900: UDP, length 106
^C
14 packets captured
14 packets received by filter
0 packets dropped by kernel
```

Packet capture at lc2_sc1 (interface between lc2 and con_ns3 namespace)

Have only the vxlan encapsulation.

```
root@6f432419fae6:/#
root@6f432419fae6:/# tcpdump -i lc2_sc2 -nn
tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
listening on lc2_sc2, link-type EN10MB (Ethernet), capture size 262144 bytes
23:31:39.905200 IP 26.0.0.2 > 28.0.0.2: GREv0, length 138: IP 21.0.0.1.40809 > 23.0.0.1.4900: UDP, length 106
23:31:39.905316 IP 28.0.0.2 > 26.0.0.2: GREv0, length 138: IP 23.0.0.1.40809 > 21.0.0.1.4900: UDP, length 106
23:31:40.929100 IP 26.0.0.2 > 28.0.0.2: GREv0, length 138: IP 21.0.0.1.40809 > 23.0.0.1.4900: UDP, length 106
23:31:40.929178 IP 28.0.0.2 > 26.0.0.2: GREv0, length 138: IP 23.0.0.1.40809 > 21.0.0.1.4900: UDP, length 106
23:31:41.953239 IP 26.0.0.2 > 28.0.0.2: GREv0, length 138: IP 21.0.0.1.40809 > 23.0.0.1.4900: UDP, length 106
23:31:41.953355 IP 28.0.0.2 > 26.0.0.2: GREv0, length 138: IP 23.0.0.1.40809 > 21.0.0.1.4900: UDP, length 106
23:31:42.977241 IP 26.0.0.2 > 28.0.0.2: GREv0, length 138: IP 21.0.0.1.40809 > 23.0.0.1.4900: UDP, length 106
23:31:42.977368 IP 28.0.0.2 > 26.0.0.2: GREv0, length 138: IP 23.0.0.1.40809 > 21.0.0.1.4900: UDP, length 106
23:31:44.001246 IP 26.0.0.2 > 28.0.0.2: GREv0, length 138: IP 21.0.0.1.40809 > 23.0.0.1.4900: UDP, length 106
23:31:44.001373 IP 28.0.0.2 > 26.0.0.2: GREv0, length 138: IP 23.0.0.1.40809 > 21.0.0.1.4900: UDP, length 106
23:31:45.025214 IP 26.0.0.2 > 28.0.0.2: GREv0, length 138: IP 21.0.0.1.40809 > 23.0.0.1.4900: UDP, length 106
23:31:45.025337 IP 28.0.0.2 > 26.0.0.2: GREv0, length 138: IP 23.0.0.1.40809 > 21.0.0.1.4900: UDP, length 106
23:31:46.049215 IP 26.0.0.2 > 28.0.0.2: GREv0, length 138: IP 21.0.0.1.40809 > 23.0.0.1.4900: UDP, length 106
23:31:46.049336 IP 28.0.0.2 > 26.0.0.2: GREv0, length 138: IP 23.0.0.1.40809 > 21.0.0.1.4900: UDP, length 106
23:31:47.073210 IP 26.0.0.2 > 28.0.0.2: GREv0, length 138: IP 21.0.0.1.40809 > 23.0.0.1.4900: UDP, length 106
23:31:47.073328 IP 28.0.0.2 > 26.0.0.2: GREv0, length 138: IP 23.0.0.1.40809 > 21.0.0.1.4900: UDP, length 106
23:31:48.097168 IP 26.0.0.2 > 28.0.0.2: GREv0, length 138: IP 21.0.0.1.40809 > 23.0.0.1.4900: UDP, length 106
23:31:48.097256 IP 28.0.0.2 > 26.0.0.2: GREv0, length 138: IP 23.0.0.1.40809 > 21.0.0.1.4900: UDP, length 106
^C
18 packets captured
18 packets received by filter
0 packets dropped by kernel
```

Packet capture at cs5:

Both VXLAN and GRE headers have been decapsulated at this point.

```

valid_ttl forever preferred_ttl forever
root@b80ed10ee269:/# tcpdump -i cs5_br3
tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
listening on cs5_br3, link-type EN10MB (Ethernet), capture size 262144 bytes
23:14:14.305194 IP 20.0.0.1 > 20.0.0.3: ICMP echo request, id 84, seq 26, length 64
23:14:14.305225 IP 20.0.0.3 > 20.0.0.1: ICMP echo reply, id 84, seq 26, length 64
23:14:15.329341 IP 20.0.0.1 > 20.0.0.3: ICMP echo request, id 84, seq 27, length 64
23:14:15.329365 IP 20.0.0.3 > 20.0.0.1: ICMP echo reply, id 84, seq 27, length 64
23:14:16.353270 IP 20.0.0.1 > 20.0.0.3: ICMP echo request, id 84, seq 28, length 64
23:14:16.353312 IP 20.0.0.3 > 20.0.0.1: ICMP echo reply, id 84, seq 28, length 64
23:14:17.377288 IP 20.0.0.1 > 20.0.0.3: ICMP echo request, id 84, seq 29, length 64
23:14:17.377325 IP 20.0.0.3 > 20.0.0.1: ICMP echo reply, id 84, seq 29, length 64
23:14:18.401209 IP 20.0.0.1 > 20.0.0.3: ICMP echo request, id 84, seq 30, length 64
23:14:18.401234 IP 20.0.0.3 > 20.0.0.1: ICMP echo reply, id 84, seq 30, length 64
23:14:19.425281 IP 20.0.0.1 > 20.0.0.3: ICMP echo request, id 84, seq 31, length 64
23:14:19.425313 IP 20.0.0.3 > 20.0.0.1: ICMP echo reply, id 84, seq 31, length 64
23:14:20.449162 IP 20.0.0.1 > 20.0.0.3: ICMP echo request, id 84, seq 32, length 64
23:14:20.449185 IP 20.0.0.3 > 20.0.0.1: ICMP echo reply, id 84, seq 32, length 64
23:14:20.513130 ARP, Request who-has 20.0.0.3 tell 20.0.0.1, length 28
23:14:20.513145 ARP, Reply 20.0.0.3 is-at ea:36:90:91:bd:87 (oui Unknown), length 28
23:14:21.473201 IP 20.0.0.1 > 20.0.0.3: ICMP echo request, id 84, seq 33, length 64
23:14:21.473233 IP 20.0.0.3 > 20.0.0.1: ICMP echo reply, id 84, seq 33, length 64
^C
18 packets captured
18 packets received by filter
0 packets dropped by kernel

```

Pinging between the cs1 (20.0.0.1) cs3 (30.0.0.1) and cs4(30.0.0.2)

Capture at cs1 :

```

172.17.0.0 0.0.0.0 255.255.0.0 U 0 0 0 eth0
root@cla6f8c932cd:/# ping 30.0.0.1
PING 30.0.0.1 (30.0.0.1) 56(84) bytes of data.
54 bytes from 30.0.0.1: icmp_seq=1 ttl=62 time=0.280 ms
54 bytes from 30.0.0.1: icmp_seq=2 ttl=62 time=0.185 ms
54 bytes from 30.0.0.1: icmp_seq=3 ttl=62 time=0.244 ms
54 bytes from 30.0.0.1: icmp_seq=4 ttl=62 time=0.166 ms
54 bytes from 30.0.0.1: icmp_seq=5 ttl=62 time=0.168 ms
54 bytes from 30.0.0.1: icmp_seq=6 ttl=62 time=0.166 ms
54 bytes from 30.0.0.1: icmp_seq=7 ttl=62 time=0.213 ms
54 bytes from 30.0.0.1: icmp_seq=8 ttl=62 time=0.165 ms
54 bytes from 30.0.0.1: icmp_seq=9 ttl=62 time=0.166 ms
54 bytes from 30.0.0.1: icmp_seq=10 ttl=62 time=0.274 ms
^C
--- 30.0.0.1 ping statistics ---
10 packets transmitted, 10 received, 0% packet loss, time 9162ms
rtt min/avg/max/mdev = 0.165/0.202/0.280/0.047 ms
root@cla6f8c932cd:/# ping 30.0.0.2
PING 30.0.0.2 (30.0.0.2) 56(84) bytes of data.
54 bytes from 30.0.0.2: icmp_seq=1 ttl=62 time=0.283 ms
54 bytes from 30.0.0.2: icmp_seq=2 ttl=62 time=0.258 ms
54 bytes from 30.0.0.2: icmp_seq=3 ttl=62 time=0.266 ms
54 bytes from 30.0.0.2: icmp_seq=4 ttl=62 time=0.263 ms
54 bytes from 30.0.0.2: icmp_seq=5 ttl=62 time=0.220 ms
54 bytes from 30.0.0.2: icmp_seq=6 ttl=62 time=0.179 ms
54 bytes from 30.0.0.2: icmp_seq=7 ttl=62 time=0.259 ms
^C
--- 30.0.0.2 ping statistics ---
7 packets transmitted, 7 received, 0% packet loss, time 6133ms
rtt min/avg/max/mdev = 0.179/0.246/0.283/0.038 ms
root@cla6f8c932cd:/#

```

Capture at gretun-cont in lc1:

```

root@8015c76fb745:/# tcpdump -i gretun-cont -nn
bash: tcpdump: command not found
root@8015c76fb745:/# tcpdump -i gretun-cont -nn
tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
listening on gretun-cont, link-type LINUX_SLL (Linux cooked), capture size 262144 bytes
02:06:09.121326 IP 20.0.0.1 > 30.0.0.2: ICMP echo request, id 167, seq 48, length 64
02:06:09.121492 IP 30.0.0.2 > 20.0.0.1: ICMP echo reply, id 167, seq 48, length 64
02:06:10.145195 IP 20.0.0.1 > 30.0.0.2: ICMP echo request, id 167, seq 49, length 64
02:06:10.145346 IP 30.0.0.2 > 20.0.0.1: ICMP echo reply, id 167, seq 49, length 64
02:06:11.169316 IP 20.0.0.1 > 30.0.0.2: ICMP echo request, id 167, seq 50, length 64
02:06:11.180379 IP 30.0.0.2 > 20.0.0.1: ICMP echo reply, id 167, seq 50, length 64
02:06:12.169643 IP 20.0.0.1 > 30.0.0.2: ICMP echo request, id 167, seq 51, length 64
02:06:12.169745 IP 30.0.0.2 > 20.0.0.1: ICMP echo reply, id 167, seq 51, length 64
02:06:13.185133 IP 20.0.0.1 > 30.0.0.2: ICMP echo request, id 167, seq 52, length 64
02:06:13.185326 IP 30.0.0.2 > 20.0.0.1: ICMP echo reply, id 167, seq 52, length 64
^C
10 packets captured
10 packets received by filter
0 packets dropped by kernel

```

Capture at gretun in lc2:

```
root@8015c76fb745:/# tcpdump -i gretun-cont -nn
bash: tcpdump: command not found
root@8015c76fb745:/# tcpdump -i gretun-cont -nn
tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
listening on gretun-cont, link-type LINUX_SLL (Linux cooked), capture size 262144 bytes
02:06:09.121326 IP 20.0.0.1 > 30.0.0.2: ICMP echo request, id 167, seq 48, length 64
02:06:09.121492 IP 30.0.0.2 > 20.0.0.1: ICMP echo reply, id 167, seq 48, length 64
02:06:10.145195 IP 20.0.0.1 > 30.0.0.2: ICMP echo request, id 167, seq 49, length 64
02:06:10.145346 IP 30.0.0.2 > 20.0.0.1: ICMP echo reply, id 167, seq 49, length 64
02:06:11.169316 IP 20.0.0.1 > 30.0.0.2: ICMP echo request, id 167, seq 50, length 64
02:06:11.180379 IP 30.0.0.2 > 20.0.0.1: ICMP echo reply, id 167, seq 50, length 64
02:06:12.169643 IP 20.0.0.1 > 30.0.0.2: ICMP echo request, id 167, seq 51, length 64
02:06:12.169745 IP 30.0.0.2 > 20.0.0.1: ICMP echo reply, id 167, seq 51, length 64
02:06:13.185133 IP 20.0.0.1 > 30.0.0.2: ICMP echo request, id 167, seq 52, length 64
02:06:13.185326 IP 30.0.0.2 > 20.0.0.1: ICMP echo reply, id 167, seq 52, length 64
^C
10 packets captured
10 packets received by filter
```

Capture at sc2-lc1 node (connection between lc1-sc2):

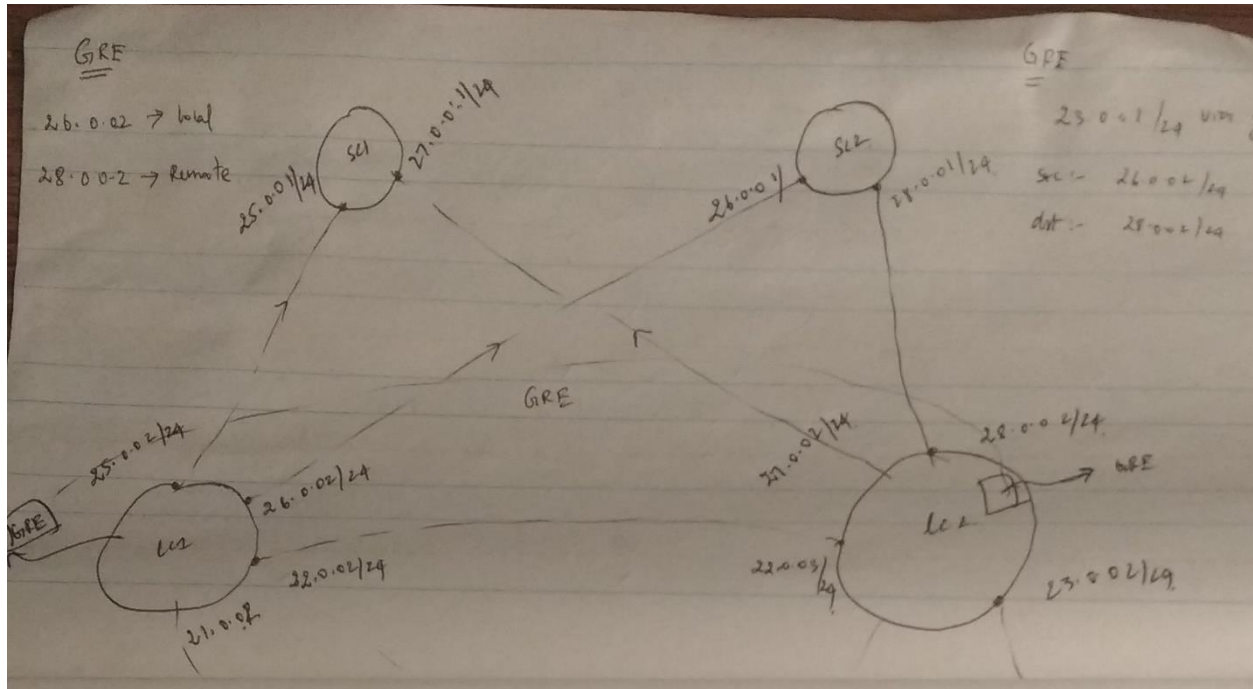
```
root@7119e645b7cb:/# tcpdump -i sc2_lc1 -nn
tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
listening on sc2_lc1, link-type EN10MB (Ethernet), capture size 262144 bytes
02:10:26.017274 IP 26.0.0.2 > 28.0.0.2: GREv0, length 88: IP 20.0.0.1 > 30.0.0.2: ICMP echo request, id 167, seq 299, length 64
02:10:26.017390 IP 28.0.0.2 > 26.0.0.2: GREv0, length 88: IP 30.0.0.2 > 20.0.0.1: ICMP echo reply, id 167, seq 299, length 64
02:10:27.041312 IP 26.0.0.2 > 28.0.0.2: GREv0, length 88: IP 20.0.0.1 > 30.0.0.2: ICMP echo request, id 167, seq 300, length 64
02:10:27.041450 IP 28.0.0.2 > 26.0.0.2: GREv0, length 88: IP 30.0.0.2 > 20.0.0.1: ICMP echo reply, id 167, seq 300, length 64
02:10:28.065124 IP 26.0.0.2 > 28.0.0.2: GREv0, length 88: IP 20.0.0.1 > 30.0.0.2: ICMP echo request, id 167, seq 301, length 64
02:10:28.065220 IP 28.0.0.2 > 26.0.0.2: GREv0, length 88: IP 30.0.0.2 > 20.0.0.1: ICMP echo reply, id 167, seq 301, length 64
02:10:29.089135 IP 26.0.0.2 > 28.0.0.2: GREv0, length 88: IP 20.0.0.1 > 30.0.0.2: ICMP echo request, id 167, seq 302, length 64
02:10:29.089223 IP 28.0.0.2 > 26.0.0.2: GREv0, length 88: IP 30.0.0.2 > 20.0.0.1: ICMP echo reply, id 167, seq 302, length 64
02:10:30.113163 IP 26.0.0.2 > 28.0.0.2: GREv0, length 88: IP 20.0.0.1 > 30.0.0.2: ICMP echo request, id 167, seq 303, length 64
02:10:30.113287 IP 28.0.0.2 > 26.0.0.2: GREv0, length 88: IP 30.0.0.2 > 20.0.0.1: ICMP echo reply, id 167, seq 303, length 64
02:10:31.137165 IP 26.0.0.2 > 28.0.0.2: GREv0, length 88: IP 20.0.0.1 > 30.0.0.2: ICMP echo request, id 167, seq 304, length 64
02:10:31.137266 IP 28.0.0.2 > 26.0.0.2: GREv0, length 88: IP 30.0.0.2 > 20.0.0.1: ICMP echo reply, id 167, seq 304, length 64
02:10:32.161115 IP 26.0.0.2 > 28.0.0.2: GREv0, length 88: IP 20.0.0.1 > 30.0.0.2: ICMP echo request, id 167, seq 305, length 64
02:10:32.161230 IP 28.0.0.2 > 26.0.0.2: GREv0, length 88: IP 30.0.0.2 > 20.0.0.1: ICMP echo reply, id 167, seq 305, length 64
^C
14 packets captured
14 packets received by filter
0 packets dropped by kernel
```

Capture at cs4 :

```
root@2d22bb398409:/# tcpdump -i cs4-br
tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
listening on cs4-br, link-type EN10MB (Ethernet), capture size 262144 bytes
02:12:07.393162 IP 20.0.0.1 > 30.0.0.2: ICMP echo request, id 167, seq 398, length 64
02:12:07.393189 IP 30.0.0.2 > 20.0.0.1: ICMP echo reply, id 167, seq 398, length 64
02:12:08.417199 IP 20.0.0.1 > 30.0.0.2: ICMP echo request, id 167, seq 399, length 64
02:12:08.417235 IP 30.0.0.2 > 20.0.0.1: ICMP echo reply, id 167, seq 399, length 64
02:12:09.441221 IP 20.0.0.1 > 30.0.0.2: ICMP echo request, id 167, seq 400, length 64
02:12:09.441266 IP 30.0.0.2 > 20.0.0.1: ICMP echo reply, id 167, seq 400, length 64
02:12:09.569031 ARP, Request who-has 30.0.0.2 tell 30.0.0.3, length 28
02:12:09.569062 ARP, Reply 30.0.0.2 is-at ae:13:f7:0e:bd:c8 (oui Unknown), length 28
02:12:10.465212 IP 20.0.0.1 > 30.0.0.2: ICMP echo request, id 167, seq 401, length 64
02:12:10.465248 IP 30.0.0.2 > 20.0.0.1: ICMP echo reply, id 167, seq 401, length 64
02:12:11.489174 IP 20.0.0.1 > 30.0.0.2: ICMP echo request, id 167, seq 402, length 64
02:12:11.489201 IP 30.0.0.2 > 20.0.0.1: ICMP echo reply, id 167, seq 402, length 64
^C
12 packets captured
12 packets received by filter
0 packets dropped by kernel
root@2d22bb398409:/#
```

(Part2):

The lc1 and lc2 and sc1 and sc2 topology and configurations are below.



Please refer to Readme file.



2.1 Iperf performance on the virtual machines:

Iperf output on the client VM:

```
iperf Done.
[root@client ~]# iperf3 -c 10.0.0.2 -u -b 0
Connecting to host 10.0.0.2, port 5201
[ 4] local 10.0.0.1 port 53014 connected to 10.0.0.2 port 5201
[ ID] Interval           Transfer     Bandwidth   Total Datagrams
[ 4]  0.00-1.01   sec  1.39 MBytes  11.6 Mbits/sec  1010
[ 4]  1.01-2.01   sec  1.37 MBytes  11.4 Mbits/sec  990
[ 4]  2.01-3.01   sec  1.10 MBytes  9.28 Mbits/sec  800
[ 4]  3.01-4.01   sec  1.08 MBytes  9.05 Mbits/sec  780
[ 4]  4.01-5.02   sec   594 KBytes  4.83 Mbits/sec  420
[ 4]  5.02-6.00   sec  1004 KBytes  8.35 Mbits/sec  710
[ 4]  6.00-7.00   sec  1.35 MBytes  11.4 Mbits/sec  980
[ 4]  7.00-8.01   sec  1.44 MBytes  12.0 Mbits/sec  1040
[ 4]  8.01-9.01   sec  1.20 MBytes  10.0 Mbits/sec  870
[ 4]  9.01-10.01  sec  1.41 MBytes  11.8 Mbits/sec  1020
- - - - -
[ ID] Interval           Transfer     Bandwidth   Jitter    Lost/Total Datagrams
[ 4]  0.00-10.01  sec  11.9 MBytes  9.97 Mbits/sec  2.512 ms  0/8619 (0%)
[ 4] Sent 8619 datagrams
```

Iperf output on the server VM:

```
[ ID] Interval      Transfer    Bandwidth      Jitter    Lost/Total Datagrams
[ 5]  0.00-10.05  sec  0.00 Bytes  0.00 bits/sec  3.904 ms  0/6645 (0%)
-----
Server listening on 5201
-----
Accepted connection from 10.0.0.1, port 58966
[ 5] local 10.0.0.2 port 5201 connected to 10.0.0.1 port 53014
[ ID] Interval      Transfer    Bandwidth      Jitter    Lost/Total Datagrams
[ 5]  0.00-1.00   sec  1.35 MBytes  11.3 Mbits/sec  1.537 ms  0/975 (0%)
[ 5]  1.00-2.02   sec  1.35 MBytes  11.1 Mbits/sec  2.793 ms  0/974 (0%)
[ 5]  2.02-3.00   sec  1.10 MBytes  9.37 Mbits/sec  1.646 ms  0/798 (0%)
[ 5]  3.00-4.00   sec  1.11 MBytes  9.35 Mbits/sec  2.621 ms  0/806 (0%)
[ 5]  4.00-5.00   sec  587 KBytes  4.81 Mbits/sec  2.886 ms  0/415 (0%)
[ 5]  5.00-6.00   sec  952 KBytes  7.79 Mbits/sec  1.419 ms  0/673 (0%)
[ 5]  6.00-7.00   sec  1.37 MBytes  11.5 Mbits/sec  1.466 ms  0/993 (0%)
[ 5]  7.00-8.00   sec  1.43 MBytes  12.0 Mbits/sec  1.340 ms  0/1032 (0%)
[ 5]  8.00-9.00   sec  1.22 MBytes  10.2 Mbits/sec  2.774 ms  0/884 (0%)
[ 5]  9.00-10.00  sec  1.44 MBytes  12.1 Mbits/sec  2.591 ms  0/1040 (0%)
[ 5] 10.00-10.06  sec  41.0 KBytes  5.23 Mbits/sec  2.512 ms  0/29 (0%)
-----
[ ID] Interval      Transfer    Bandwidth      Jitter    Lost/Total Datagrams
[ 5]  0.00-10.06  sec  0.00 Bytes  0.00 bits/sec  2.512 ms  0/8619 (0%)
-----
Server listening on 5201
-----
```

Trace output for the the server vm `ip_rcv()` function:

Ran the following commands before taking the trace output :

```
echo function_graph > current_tracer
echo ip-* > set_fttrace_filter
echo <PID of IPERF3> > set_fttrace_pid
echo > trace
ps -ef | grep iperf3 ----- to find pid of iperf
```

```
0)          ip_local_deliver() {
0) 0.857 us   |   ip_local_deliver_finish();
0) 2.537 us   |   }
0) 3.707 us   |   }
0) 6.843 us   |   }
0)          ip_rcv() {
0)          ip_rcv_finish() {
0)          ip_local_deliver() {
0)          ip_local_deliver_finish();
0) 4.065 us   |   }
0) ! 6563.898 us |   }
0) ! 6564.891 us |   }
0) ! 6567.814 us |   }
0)          ip_rcv() {
0)          ip_rcv_finish() {
0)          ip_local_deliver() {
0)          ip_local_deliver_finish();
0) 1.414 us   |   }
0) 5.236 us   |   }
0) 6.797 us   |   }
0) + 15.932 us |   }
0)          ip_rcv() {
0)          ip_rcv_finish() {
0)          ip_local_deliver() {
0)          ip_local_deliver_finish();
0) 1.427 us   |   }
0) 3.436 us   |   }
0) 4.818 us   |   }
0) 9.405 us   |   }
0)          ip_rcv() {
```

Here the `ip_rcv` function is taking around 6000 us, because of which the throughput is low when compared to the containers. 

2.2 Iperf throughput on the container side :

Iperf calculation on client container:

```
iperf Done.
root@14dfdf86ed2d:/# iperf3 -c 172.17.0.3 -u -b 0
Connecting to host 172.17.0.3, port 5201
[ 4] local 172.17.0.4 port 50018 connected to 172.17.0.3 port 5201
[ ID] Interval      Transfer    Bandwidth  Total Datagrams
[ 4] 0.00-1.00 sec    533 MBytes  4.47 Gbits/sec  68180
[ 4] 1.00-2.00 sec    562 MBytes  4.71 Gbits/sec  71920
[ 4] 2.00-3.00 sec    546 MBytes  4.58 Gbits/sec  69930
[ 4] 3.00-4.00 sec    534 MBytes  4.48 Gbits/sec  68300
[ 4] 4.00-5.00 sec    461 MBytes  3.87 Gbits/sec  59010
[ 4] 5.00-6.00 sec    614 MBytes  5.15 Gbits/sec  78590
[ 4] 6.00-7.00 sec    602 MBytes  5.05 Gbits/sec  77010
[ 4] 7.00-8.00 sec    642 MBytes  5.38 Gbits/sec  82170
[ 4] 8.00-9.00 sec    524 MBytes  4.40 Gbits/sec  67120
[ 4] 9.00-10.00 sec   535 MBytes  4.49 Gbits/sec  68530
-----
[ ID] Interval      Transfer    Bandwidth  Jitter    Lost/Total Datagrams
[ 4] 0.00-10.00 sec  5.42 GBytes  4.66 Gbits/sec  0.006 ms  131428/710760 (18%)
[ 4] Sent 710760 datagrams
```

Iperf calculation on server container:

```
[ 5] 0.00-10.00 sec  0.00 Bytes  0.00 Gbits/sec  0.004 ms  369119/674330 (13%)
-----
Server listening on 5201
Accepted connection from 172.17.0.4, port 59192
[ 5] local 172.17.0.3 port 5201 connected to 172.17.0.4 port 50018
[ ID] Interval      Transfer    Bandwidth  Jitter    Lost/Total Datagrams
[ 5] 0.00-1.00 sec    348 MBytes  2.92 Gbits/sec  0.057 ms  20056/64594 (31%)
[ 5] 1.00-2.00 sec    486 MBytes  4.07 Gbits/sec  0.011 ms  11916/74071 (16%)
[ 5] 2.00-3.00 sec    491 MBytes  4.12 Gbits/sec  0.019 ms  5440/68277 (8%)
[ 5] 3.00-4.00 sec    452 MBytes  3.79 Gbits/sec  0.005 ms  9567/67425 (14%)
[ 5] 4.00-5.00 sec    363 MBytes  3.04 Gbits/sec  0.114 ms  12088/58570 (21%)
[ 5] 5.00-6.00 sec    447 MBytes  3.75 Gbits/sec  0.003 ms  21803/79008 (28%)
[ 5] 6.00-7.00 sec    462 MBytes  3.88 Gbits/sec  0.005 ms  19435/78570 (25%)
[ 5] 7.00-8.00 sec    604 MBytes  5.07 Gbits/sec  0.059 ms  3506/80842 (4.3%)
[ 5] 8.00-9.00 sec    387 MBytes  3.24 Gbits/sec  0.004 ms  19944/69439 (29%)
[ 5] 9.00-10.00 sec   459 MBytes  3.85 Gbits/sec  0.006 ms  7212/65995 (11%)
[ 5] 10.00-10.04 sec   27.4 MBytes  5.32 Gbits/sec  0.006 ms  461/3969 (12%)
-----
[ ID] Interval      Transfer    Bandwidth  Jitter    Lost/Total Datagrams
[ 5] 0.00-10.04 sec  0.00 Bytes  0.00 bits/sec  0.006 ms  131428/710760 (18%)
-----
Server listening on 5201
```

Trace output for the the server container ip_rcv() function:

```
echo function_graph > current_tracer
echo ip-* > set_fttrace_filter
echo <PID of IPERF3> > set_fttrace_pid
echo > trace
ps -ef | grep iperf3 ----- to find pid of iperf
```

```
1) 9.031 us      }
1) + 16.851 us  }
1) + 24.270 us  }
1)              ip_rcv() {
1) 0.080 us      ip_sabotage_in [br_netfilter()];
1)              ip_rcv_finish() {
1)              ip_local_deliver() {
1) + 10.970 us    ip_local_deliver_finish();
1) + 13.745 us    }
1) + 16.344 us    }
1) + 23.949 us    }
1) 0.272 us      ip_mc_drop_socket();
1)              ip_queue_xmit() {
1)              ip_local_out() {
```

Here the ip_rcv function is taking around 10 us because of which the throughput is higher which is around 4gbps.

2.3 comparison between the iperf performance in containers & Virtual machines :

From the trace output :



ip_rcv() is taking around 6ms in the virtual machines where as ip_rcv() function is taking in around 10ms

Which is the reason the container iperf throughput is higher in containers when compared to iperf throughput in VM

Trace output in containers:

```
1) 9.031 us      }
1) + 16.851 us  }
1) + 24.270 us  }
1)              ip_rcv() {
1) 0.080 us      ip_sabotage_in [br_netfilter]();
1)              ip_rcv_finish() {
1)              ip_local_deliver() {
1)              ip_local_deliver_finish();
1)              }
1)              }
1) + 10.970 us   }
1) + 13.745 us   }
1) + 16.344 us   }
1) + 23.949 us   }
1) 0.272 us      ip_mc_drop_socket();
1)              ip_queue_xmit() {
1)              ip_local_out() {
```

Trace output in Virtual machines:

```
0)              ip_local_deliver() {
0) 0.857 us      ip_local_deliver_finish();
0) 2.537 us      }
0) 3.707 us      }
0) 6.843 us      }
0)              ip_rcv() {
0)              ip_rcv_finish() {
0)              ip_local_deliver() {
0)              ip_local_deliver_finish();
0) 4.065 us      }
0) ! 6563.898 us }
0) ! 6564.891 us }
0) ! 6567.814 us }
0)              ip_rcv() {
0)              ip_rcv_finish() {
0)              ip_local_deliver() {
0)              ip_local_deliver_finish();
0) 1.414 us      }
0) 5.236 us      }
0) 6.797 us      }
0) + 15.932 us   }
0)              ip_rcv() {
0)              ip_rcv_finish() {
0)              ip_local_deliver() {
0)              ip_local_deliver_finish();
0) 1.427 us      }
0) 3.436 us      }
0) 4.818 us      }
0) 9.405 us      }
0)              ip_rcv() {
```

2.4 iperf performance with different packet sizes in container:

Iperf performance with packet size of 200:

```
iperf Done.
root@14d4df86ed2d:/# iperf3 -c 172.17.0.3 -u -b 0 -l 200
Connecting to host 172.17.0.3, port 5201
[ 4] local 172.17.0.4 port 43527 connected to 172.17.0.3 port 5201
[ ID] Interval      Transfer       Bandwidth     Total Datagrams
[ 4] 0.00-1.00    sec 14.4 MBytes  121 Mbits/sec  75650
[ 4] 1.00-2.00    sec 13.7 MBytes  115 Mbits/sec  71940
[ 4] 2.00-3.00    sec 16.3 MBytes  137 Mbits/sec  85500
[ 4] 3.00-4.00    sec 16.2 MBytes  136 Mbits/sec  85140
[ 4] 4.00-5.00    sec 13.1 MBytes  110 Mbits/sec  68750
[ 4] 5.00-6.00    sec 13.2 MBytes  110 Mbits/sec  69050
[ 4] 6.00-7.00    sec 17.2 MBytes  144 Mbits/sec  90190
[ 4] 7.00-8.00    sec 13.4 MBytes  112 Mbits/sec  70350
[ 4] 8.00-9.00    sec 14.8 MBytes  124 Mbits/sec  77660
[ 4] 9.00-10.00   sec 16.7 MBytes  140 Mbits/sec  87700
-- -- -- -- --
[ ID] Interval      Transfer       Bandwidth     Jitter      Lost/Total Datagrams
```

Iperf performance with packet size of 400:

```
iperf Done.
root@14d4df86ed2d:/# iperf3 -c 172.17.0.3 -u -b 0 -l 400
Connecting to host 172.17.0.3, port 5201
[ 4] local 172.17.0.4 port 58897 connected to 172.17.0.3 port 5201
[ ID] Interval      Transfer       Bandwidth     Total Datagrams
[ 4] 0.00-1.00    sec 15.7 MBytes  132 Mbits/sec  41190
[ 4] 1.00-2.00    sec 26.9 MBytes  226 Mbits/sec  70550
[ 4] 2.00-3.00    sec 31.9 MBytes  267 Mbits/sec  83570
[ 4] 3.00-4.00    sec 34.8 MBytes  292 Mbits/sec  91200
[ 4] 4.00-5.00    sec 31.9 MBytes  267 Mbits/sec  83560
[ 4] 5.00-6.00    sec 21.9 MBytes  184 Mbits/sec  57510
[ 4] 6.00-7.00    sec 37.5 MBytes  314 Mbits/sec  98200
[ 4] 7.00-8.00    sec 32.0 MBytes  269 Mbits/sec  83960
[ 4] 8.00-9.00    sec 30.8 MBytes  258 Mbits/sec  80760
[ 4] 9.00-10.00   sec 28.8 MBytes  241 Mbits/sec  75460
-- -- -- -- --
[ ID] Interval      Transfer       Bandwidth     Jitter      Lost/Total Datagrams
[ 4] 0.00-10.00   sec 292 MBytes  245 Mbits/sec  0.005 ms    46891/765960 (6.1%)
[ 4] Sent 765960 datagrams

iperf Done.
root@14d4df86ed2d:/#
```

Iperf performance with packet size of 800:

```
iperf Done.
root@14d4df86ed2d:/# iperf3 -c 172.17.0.3 -u -b 0 -l 800
Connecting to host 172.17.0.3, port 5201
[ 4] local 172.17.0.4 port 44854 connected to 172.17.0.3 port 5201
[ ID] Interval      Transfer       Bandwidth     Total Datagrams
[ 4] 0.00-1.00    sec 63.9 MBytes  536 Mbits/sec  83810
[ 4] 1.00-2.00    sec 50.2 MBytes  421 Mbits/sec  65770
[ 4] 2.00-3.00    sec 55.8 MBytes  468 Mbits/sec  73120
[ 4] 3.00-4.00    sec 57.9 MBytes  486 Mbits/sec  75940
[ 4] 4.00-5.00    sec 59.0 MBytes  495 Mbits/sec  77300
[ 4] 5.00-6.00    sec 64.9 MBytes  545 Mbits/sec  85130
[ 4] 6.00-7.00    sec 65.7 MBytes  551 Mbits/sec  86060
[ 4] 7.00-8.00    sec 62.3 MBytes  523 Mbits/sec  81720
[ 4] 8.00-9.00    sec 78.7 MBytes  660 Mbits/sec  103110
[ 4] 9.00-10.00   sec 57.6 MBytes  483 Mbits/sec  75550
-- -- -- -- --
[ ID] Interval      Transfer       Bandwidth     Jitter      Lost/Total Datagrams
[ 4] 0.00-10.00   sec 616 MBytes  517 Mbits/sec  0.011 ms    99103/807510 (12%)
[ 4] Sent 807510 datagrams

iperf Done.
root@14d4df86ed2d:/#
```

With doubling the packet size the throughput is getting doubled

With increase in packet size the throughput is getting increased because

The number of cpu cycles taken to send packets is same irrespective of packet size

where as in smaller packet the number of cpu cycles taken to send the same data is higher when compared to higher sized packets.

Which is the reason the performance is better in higher sized packets.

Ip_rcv() function is taking time around 56us in case of packet size = 400:

```
) + 16.688 us      }
) + 18.917 us      }
) 0.048 us         ip_sabotage_in [br_netfilter]();
)                  ip_rcv() {
)                  ip_rcv_finish() {
)                  ip_local_deliver() {
)                  ip_local_deliver_finish();
)                  }
)                  }
)                  }
) + 48.174 us      }
) + 48.725 us      }
) + 49.577 us      }
) + 55.151 us      }
) + 55.840 us      }
)                  ip_queue_xmit() {
```

Ip_rcv() function takes time of around 28 us in case of packet size =800

```
1)                  ip_rcv() {
1) 0.176 us         ip_sabotage_in [br_netfilter]();
1)                  ip_rcv_finish() {
1)                  ip_local_deliver() {
1) + 11.806 us      ip_local_deliver_finish();
1) + 16.176 us      }
1) + 19.548 us      }
1) + 28.897 us      }
3) 0.171 us         ip_mc_drop_socket();
3)                  ip_queue_xmit() {
3)                  ip_local_out() {
3) 0.090 us         ip_send_check();
```

The time taken by the ip_rcv() function becomes half in packet size of 800 when compared to the ip_rcv() function time if packet size is 400